

Search Algorithms: Object-Oriented Implementation (Part E)

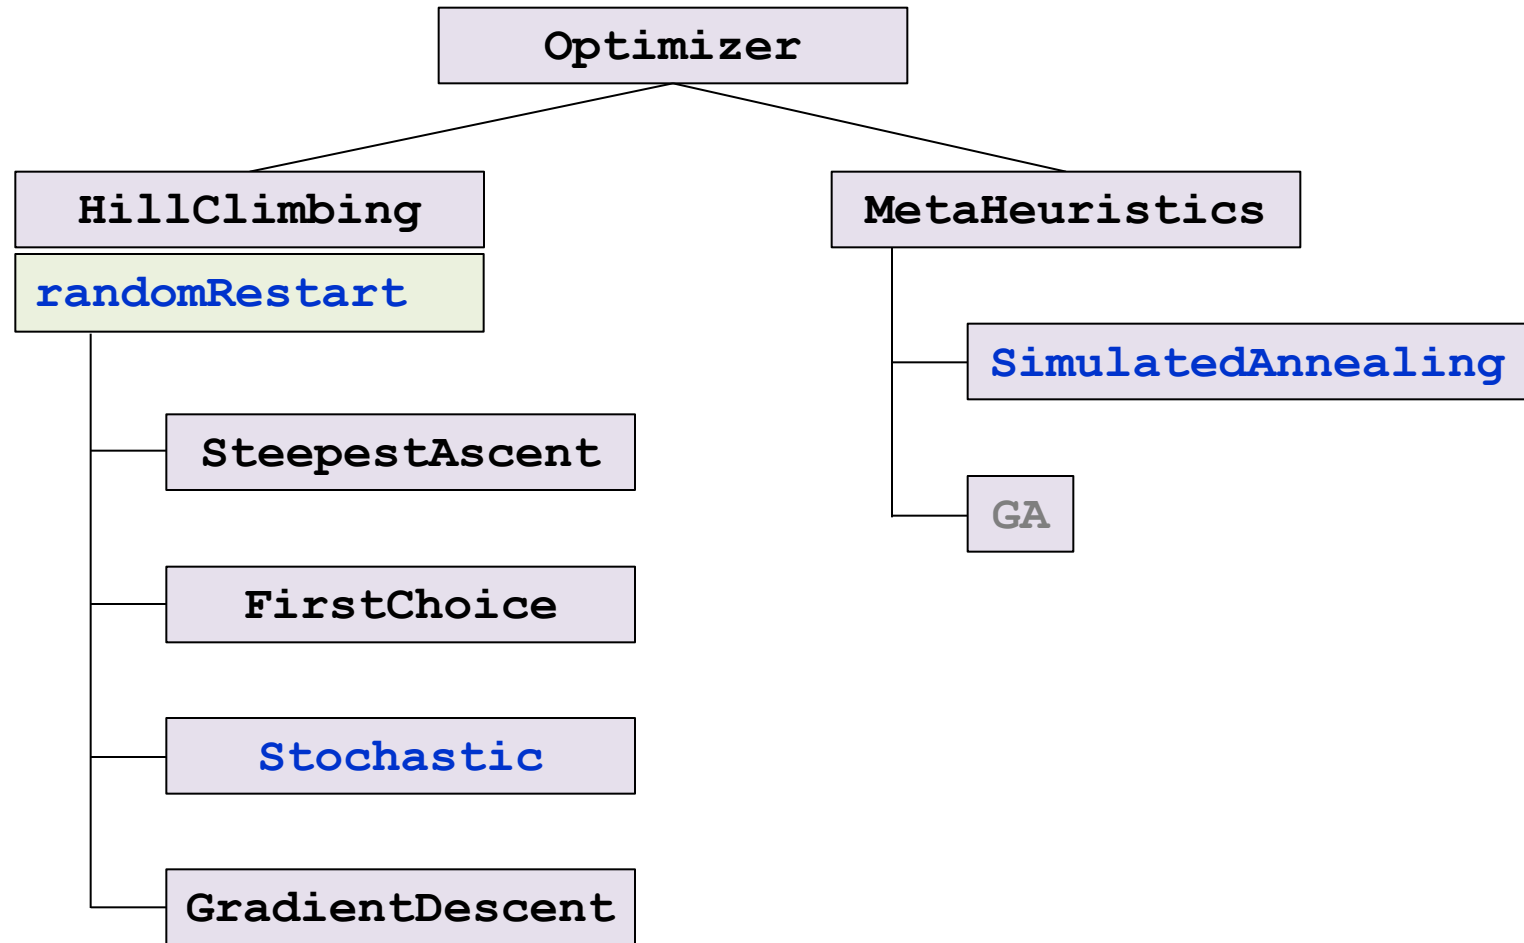
Contents

- Conventional vs. AI Algorithms
- Local Search Algorithms
- Genetic Algorithm
- Implementing Hill-Climbing Algorithms
- Defining 'Problem' Class
- Adding Gradient Descent
- Defining 'HillClimbing' Class
- **Adding More Algorithms and Classes**
- Adding Genetic Algorithm
- Experiments

Adding More Algorithms and Classes

- We add three more search algorithms: **stochastic hill climbing**, **random-restart**, and **simulated annealing**
 - This requires a more general class hierarchy to cover various algorithms
 - We define a new class **optimizer** and let it have two subclasses **HillClimbing** and **MetaHeuristics**
 - **HillClimbing** has four child classes, each for a specific hill climber
 - The random-restart algorithm becomes a method of **HillClimbing**, which can be inherited to all the child hill climbers
 - **MetaHeuristics** has two child classes, **SimulatedAnnealing** and **GA** (to be implemented later)

Adding More Algorithms and Classes



Adding More Algorithms and Classes

- Notice that `randomRestart` is a wrapper around any hill-climbing algorithm, which runs the algorithm for a given number of times
 - Each hill climber can have not only a `run` method for its own algorithm but also a `randomRestart` method as a wrapper that keeps calling the run method
 - Instead of having the same `randomRestart` method in duplicate in all the individual hill-climber subclasses, we make it a method of the parent class `HillClimbing` so that it can be inherited to all the subclasses

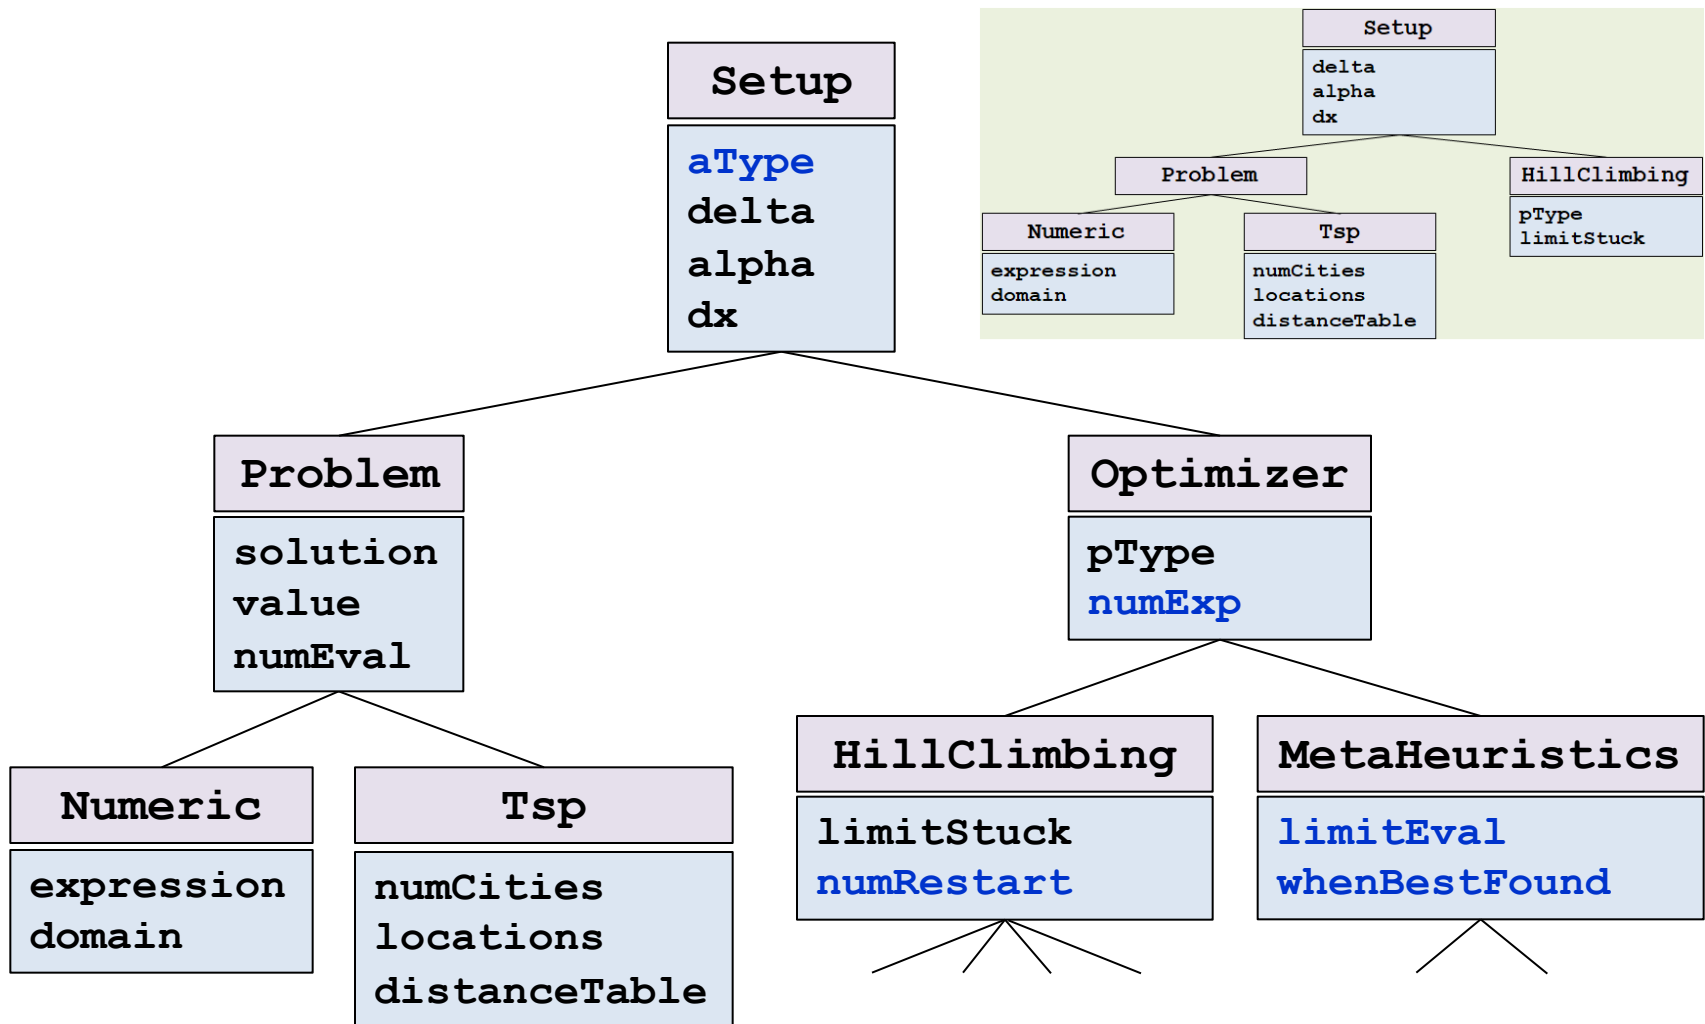
Adding More Algorithms and Classes

- When running any of the hill-climbing algorithms
 - An object `alg` is created for one of the four hill-climber subclasses chosen by the user (e.g., `alg = steepestAscent()`)
 - Then, the `randomRestart` method is called from that object (`alg.randomRestart`)
 - See the `createOptimizer` and `conductExperiment` function of the main program
 - Notice that `randomRestart` keeps calling the `self.run` method that is the algorithm of the chosen hill-climber
- In the case of metaheuristic algorithms, simply the class object of a chosen algorithm is created and run

Adding More Algorithms and Classes

- Due to the variety of algorithms and related parameters, the user interface is revised to get setup information from a file
- There are four variables newly added to `optimizer` class hierarchy
 - `numExp`: total number of experiments to be conducted (to `optimizer`)
 - `limitEval`: total number of evaluations until termination for metaheuristic algorithms (to `MetaHeuristics`)
 - `whenBestFound`: the number of iterations taken until the best solution has first been found (to `MetaHeuristics`)
 - `numRestart`: number of random restarts for the random-restart algorithm (to `HillClimbing`)

Adding More Algorithms and Classes



Adding More Algorithms and Classes

- `aType` is newly added to `Setup`
 - It is referenced by the revised `report` method of `Problem`
 - It is also referenced by `displaySetting` of `HillClimbing`
- Because of the new `optimizer` class hierarchy, `pType` has moved up from `HillClimbing` to `Optimizer` because it is referenced by the revised `displaySetting` method of `Optimizer`
 - A variable `v` is better placed at the lowest possible class `c` such that `v` is used in that class or in its subclasses

Adding More Algorithms and Classes

- Notice that an instance of an object in a class hierarchy is created for a leaf class most of the times
 - Things in the higher classes are all inherited unless otherwise specified
 - To inherit methods of the same names, it should be stated explicitly (e.g., `Optimizer.setVariables(self, parameters)`)
 - Otherwise, the methods in the upper classes will simply be overridden
 - Whenever a new instance is created, all the variables up to the base class must be set to appropriate values by calling the relevant methods such as `setVariables`

The New Main Program

- The main program is changed significantly to meet the new requirements
 - The program should be able to support multiple experiments requested by the user
 - Experimental settings and other information should be read from a file provided by the user

The New Main Program

- `main()`:
 - Reads information from a file and creates a problem `p` (`Problem` object) to be solved and an optimizer `alg` (`Optimizer` object) to be used (`readPlanAndCreate`)
 - Conducts experiments and obtains the result (`conductExperiment`)
 - Describes the problem just solved (`p.describe`)
 - Shows the settings of experiment (`alg.displayNumExp` and `alg.displaySetting`)
 - Reports the result of experiment (`p.report`)

The New Main Program

- `readPlanAndCreate()`:
 - Reads setup information from a file and stores them in the variable `parameters` (`readValidPlan`)
 - Creates `Problem` object `p` (`createProblem`)
 - Creates `optimizer` object `alg` (`createOptimizer`)
 - Returns `p` and `alg`
- `readValidPlan()`:
 - Reads setup information from a file and stores them in `parameters` (`readPlan`)
 - Keeps querying the user if gradient descent is chosen for TSP
 - Returns `parameters`

The New Main Program

- **readPlan()**:
 - Obtains a file name from the user
 - Prepares a dictionary variable **parameters** to store the information
 - Fills out **parameters** dictionary by reading the given file
 - All the information are numeric values except one file name (the file containing specifics of the target problem)
 - When reading the file, lines beginning with '#' are all skipped (**lineAfterComments**)
 - Returns **parameters**

The New Main Program

- **lineAfterComments () :**
 - Skips the lines beginning with the symbol '#' and returns the first line beginning with no '#'
- **createProblem(parameters) :**
 - Creates a **Numeric** or **Tsp** object **p** depending on the type of problem chosen
 - Sets some relevant variables of **p** in the class hierarchy with the values in **parameters** (**p.setVariables**)
 - Returns a specific problem instance **p**

The New Main Program

- `createOptimizer(parameters):`
 - Prepares a dictionary `optimizers` of algorithm class names (`optimizers`) that can be indexed by `aType`
 - Creates an object `alg` of the targeted algorithm by applying the `eval` function to the string of the name of the algorithm class (`eval(optimizers[aType])`)
 - Sets the class variables of `alg` with the values in `parameters` (`alg.setVariables`)
 - Returns `alg` as the created optimizer object

The New Main Program

- `conductExperiment(p, alg):`
 - Solves the problem `p` with the chosen optimizer `alg` and collects the result of each individual experiment
 - If the chosen algorithm is a hill climber, then the random restart algorithm is called (`alg.randomRestart`)
 - Otherwise, its `run` method is called (`alg.run`)
 - Repeats experiment multiple times (`numExp`) if requested and collects the results in a few local variables
 - Stores the final summary result (`p.storeExpResult`)

Changes to 'Problem' Class

- For recording the results of experiments, the `Problem` class is revised to have the following additional variables:
 - `pFileName`: name of the file containing problem specifics
 - `bestSolution`: best solution found in n different experiments
 - `bestMinimum`: objective value of `bestSolution`
 - `avgMinimum`: average objective value of the best solutions obtained from n experiments
 - `avgNumEval`: average number of evaluations made in n different experiments
 - `sumOfNumEval`: total number of evaluations made all through n experiments
 - `avgWhen`: average iteration when the best solution first appears in n experiments

Changes to 'Problem' Class

- Accordingly, several new methods are added to handle those variables
 - Three new accessors `getSolution`, `getValue`, `getNumEval` are necessary for `conductExperiment` of the main program to conduct multiple experiments
 - The new method `storeExpResult` is also necessary for `conductExperiment` to store the result of experiment after finishing all the experiments
- The `report` method has been revised to display the summary result of multiple experiments in an organized fashion
 - `report` in the base class prints messages that are common to both numerical optimization problem and TSP

Changes to 'Problem' Class

- Those in the subclasses print messages specific to the type of problem just solved
- The `reportNumEvals` method prints out the total number of evaluations regardless of the problem type
 - However, it does nothing when the algorithm used is simulated annealing or GA because the number of evaluations for them is predetermined
 - It is separated from `report` because we want the result messages printed out in some appropriate order when mixed together with the messages generated from the subclasses
 - Call to `reportNumEvals` is made within `report` of the subclasses at its last line

'Optimizer' Class

- `optimizer` has two variables `pType` and `numExp` to store common information about experimental settings:
 - The methods in the class hierarchy of `optimizer` are extended versions of those previously existed in `HillClimbing`
 - New accessor methods `getWhenBestFound` (in `MetaHeuristics`) and `getNumExp` (in `optimizer`) are added for being used by the `conductExperiment` function of the main program
- Both 'random.py' and 'math.py' should be imported to the 'optimizer.py' file because the methods for stochastic hill climbing and simulated annealing algorithms need them

'Optimizer' Class

- `HillClimbing` now has two variables: `limitStuck` and `numRestart`
 - `pType` has moved up to `optimizer`
- `MetaHeuristics` is intended to be a parent of `SimulatedAnnealing` and `GA`
 - There are two variables `limitEval` and `whenBestFound`
 - `displaySettings` method prints out `limitEval` (total number of evaluations until termination)

HillClimbing
<code>limitStuck</code> <code>numRestart</code>
<code>__init__</code> <code>setVariables</code> <code>displaySetting</code> <code>randomRestart</code>

MetaHeuristics
<code>limitEval</code> <code>whenBestFound</code>
<code>__init__</code> <code>setVariables</code> <code>getWhenBestFound</code> <code>displaySetting</code>

Changes to 'HillClimbing' Class

- Changes of variables:
 - Under the new `optimizer` class hierarchy, the variable `pType` is moved up from `HillClimbing` to `Optimizer`
 - A variable `numRestart` is newly added
 - The `setVariables` method is revised accordingly
- Changes to the `displaySetting` method:
 - The part of printing the mutation step size is moved up to `Optimizer`
 - Now it prints out setting information related only to the variables of its own class `HillClimbing`

Changes to 'HillClimbing' Class


- A new method `randomRestart` is added as described before
 - It keeps calling `self.run` for a given number of times (`self._numRestart`) and stores the best solution found (`p.storeResult`)

Stochastic Hill Climbing

- The algorithm of stochastic hill climbing is the same as first-choice hill climbing except the way a successor is chosen
 - First-choice hill climbing generates a single successor randomly
 - Stochastic hill climbing generates multiple neighbors and then selects one from them at random by a probability proportional to the quality
 - The algorithm is implemented as the `run` method as before

Stochastic Hill Climbing

- `stochasticBest(self, neighbors, p):`
 - Obtains a list of evaluation values of `neighbors` (`valuesForMin`)
 - Here, smaller values are better
 - Converts the list to the one in which larger values are better (`valuesForMax`)
 - Each original value is subtracted from a large enough value (the maximum plus one to avoid zero)
 - Chooses at random from `valuesForMax` with probability proportional to its value
 - Returns the chosen neighbor together with its original evaluation value



Simulated Annealing

- There is one variable `numSample` storing the number of samples used to heuristically determine an initial temperature
 - It is currently preset to 100
- `run()`:
 - Starts from a random initial point
 - While the algorithm runs for `limitEval` iterations or stops when the temperature is zero, it uses another variable `whenBestFound` to record when the best-so-far solution has first been found
 - The temperature decreases every iteration according to an annealing schedule (`self.tSchedule(t)`)
 - The initial temperature is heuristically determined so that the probability of accepting a worse neighbor becomes 0.5 initially (`self.initTemp(p)`)

Simulated Annealing

- The probability of accepting a worse neighbor is $\exp(-\Delta E / t)$, where ΔE is the difference of the evaluation values and t is the temperature
- `initTemp(self, p):`
 - Takes k (`= self._numSample`) random samples and their neighbors from the domain of problem `p`
 - Calculates the average ΔE of their differences
 - Calculates the temperature t such that $\exp(-\Delta E / t) = 0.5$, and returns t
- `tSchedule(self, t):`
 - Calculates the next temperature using a simple formula, and returns it